

Using

Binarized Neural Network

to

Compress DNN

Xianda (Bryce) Xu

xxu373@wisc.edu

November 7th

Why is model compression so important ?

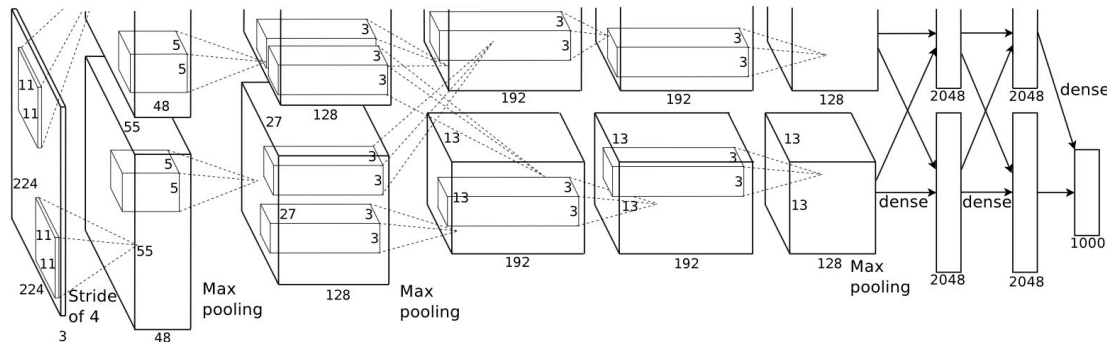


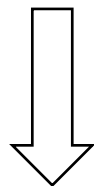
Figure 2: An illustration of the architecture of our CNN, explicitly showing the delineation of responsibilities between the two GPUs. One GPU runs the layer-parts at the top of the figure while the other runs the layer-parts at the bottom. The GPUs communicate only at certain layers. The network's input is 150,528-dimensional, and the number of neurons in the network's remaining layers is given by 253,440–186,624–64,896–64,896–43,264–4096–4096–1000.

Figure 1. AlexNet architecture

(ImageNet Large Scale Visual Recognition Challenge)

Top-1 accuracy: 57.1%

Top-5 accuracy: 80.2%



~ 60 M Parameters !

Problem 1: Computation Cost

$$A = \sigma(X \cdot W^T + B)$$

Multiplication is energy & time consuming !

Problem 2: Memory Cost

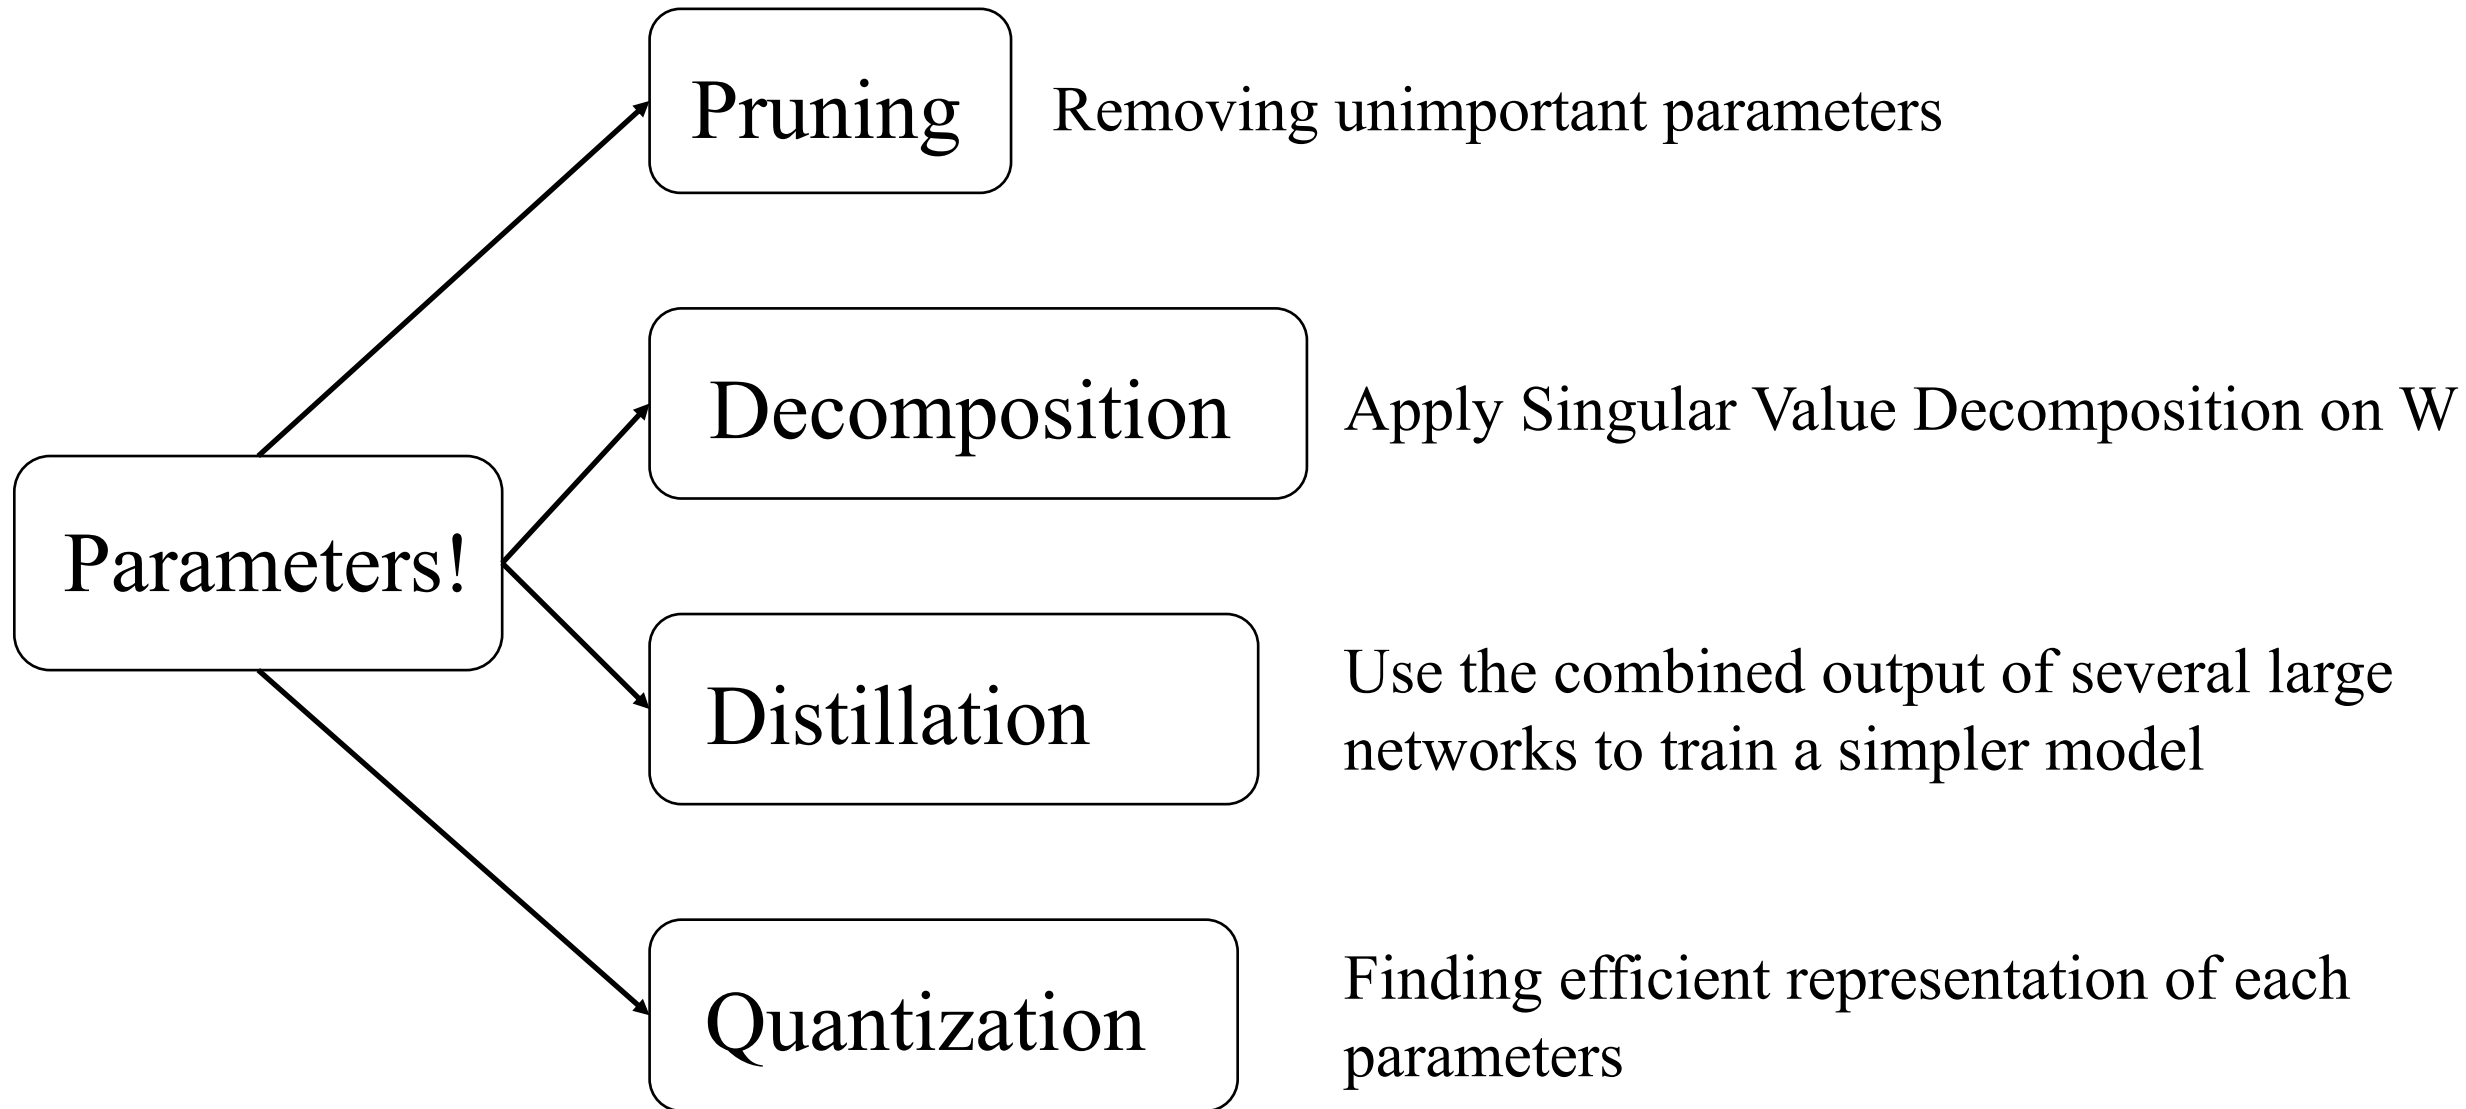
If float-32

60 M Parameters = 240 MB Memory !

However,

The energy and memory are limited on mobile devices and embedded devices !

How can we compress DNN ?



What is BNN ?

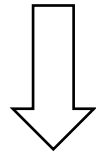
In brief, it binarizes parameters and activations to +1 and -1

Why should we choose BNN ?

-- Reduce memory cost

Full-precision parameter takes 32 bits

Binary parameter takes 1 bit

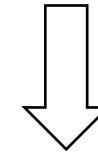


Compress network by 32-X theoretically

-- Save energy and speed up

Full-precision multiplication: \bullet

Binary multiplication: \odot (XOR)



Multiply-accumulations can be replaced
by XOR and bit-count

How do we implement BNN ?

Problem 1: How to binarize ?

-- Stochastic Binarization

$$w_b = \begin{cases} +1 & \text{with probability } p = \sigma(w), \\ -1 & \text{with probability } 1 - p. \end{cases}$$

where σ is the “hard sigmoid” function:

$$\sigma(x) = \text{clip}(\frac{x+1}{2}, 0, 1) = \max(0, \min(1, \frac{x+1}{2}))$$

-- Deterministic Binarization ☺

$$x^b = \text{Sign}(x) = \begin{cases} +1 & \text{if } x \geq 0, \\ -1 & \text{otherwise,} \end{cases}$$

Though Stochastic Binarization seems more reasonable, we prefer deterministic binarization for its efficiency.

Problem 2: When to binarize ?

-- Forward propagation

1. First Layer:

We do not binarize the input but binarize the Ws and As

2. Hidden Layers:

We binarize all the Ws and As

3. Output Layer:

We binarize Ws and only binarize the output in training

-- Back-propagation

We do not binarize gradients in back-propagation

But we have to clip weights when we update them

How do we implement BNN ?

Problem 3: How to do back-propagation ?

Recall: We calculate the gradients of the loss function L with respect to I_l , the input of the l layer.

$$g_l = \frac{\partial L}{\partial I_l}$$

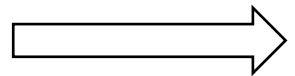
The layer activation gradients: 

$$g_{l-1} = g_l W_l^T$$

The layer weight gradients:

$$g_{W_l} = g_l I_{l-1}^T$$

But since we use Binarizing function, gradients are all zero !



Straight-Through Estimator !

Straight-Through Estimator (STE)

Yoshua Bengio etc. Estimating or propagating gradients through stochastic neurons for conditional computation (15 Aug 2013).

Adapted, for hidden layers:

For the last layer: $g_{a_L} = \frac{\partial C}{\partial a_L}$

For hidden layers: (map sign(x))

$$g_{a_k} = g_{a_k^b} 1_{|a_k| \leq 1} \quad \text{STE}$$

$$g_{s_k} = BN(g_{a_k}) \quad \text{Back Batch Norm}$$

$$g_{a_{k-1}^b} = g_{s_k} W_k^b \quad \text{Activation gradients}$$

$$g_{W_k^b} = g_{s_k}^T a_{k-1}^b \quad \text{Weight gradients}$$

STE = The gradient on Htanh(x)

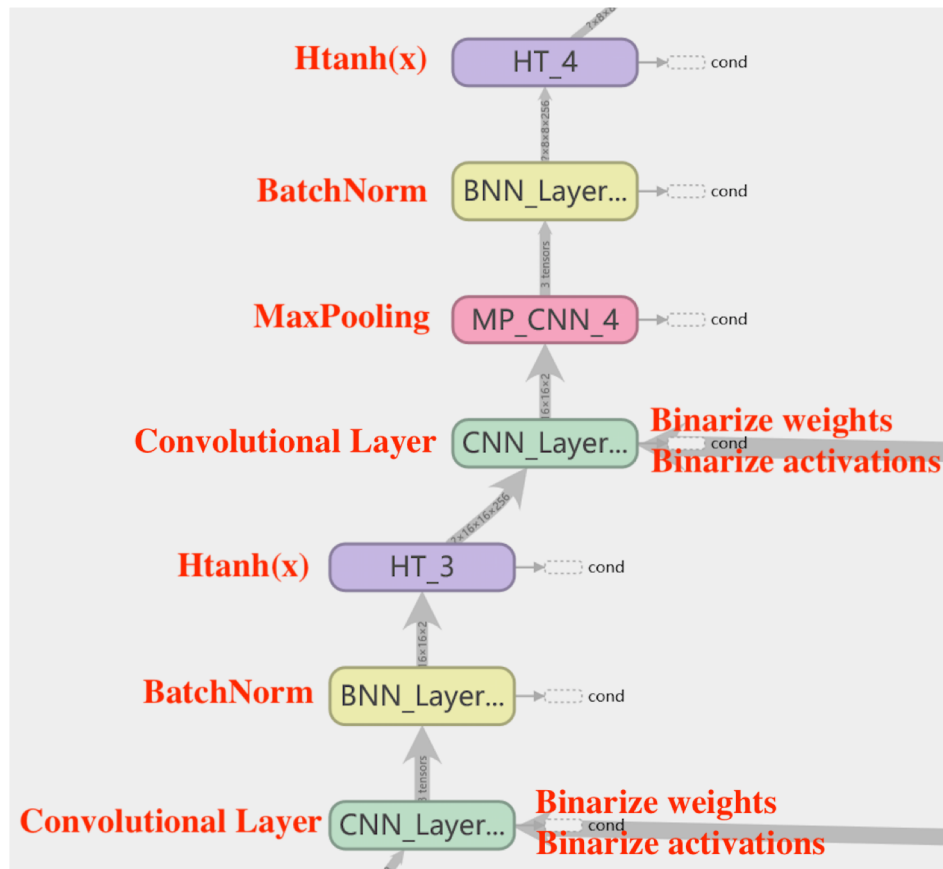
So, we use Htanh(x) as our activation function

$$H \tanh(x) = Clip(x, -1, 1)$$

How was my experiment ?

The architecture of BNN in this paper (fed by Cifar-10).

The block



In this paper:

The validation accuracy
89%

My experiment:

The training accuracy
95%

The validation accuracy
84%

```
AlexNet
Input
  Cifar-10:32x32x3
SpatialConvolution (Weight Only) 'VALID'
  32x32x3 -> 32x32x128
BatchNormalization
HardTanh
SpatialConvolution_1
  32x32x128 -> 32x32x128
MaxPooling
  32x32x128 -> 16x16x128
BatchNormalization
HardTanh
SpatialConvolution_2
  16x16x128 -> 16x16x256
BatchNormalization
HardTanh
SpatialConvolution_3
  16x16x256 -> 16x16x256
MaxPooling
  16x16x256 -> 8x8x256
BatchNormalization
HardTanh
SpatialConvolution_4
  8x8x256 -> 8x8x512
BatchNormalization
HardTanh
SpatialConvolution_5
  8x8x512 -> 8x8x512
MaxPooling
  8x8x512 -> 4x4x512
BatchNormalization
HardTanh
FC_1
  8192 -> 1024
BatchNormalization
HardTanh
FC_2
  1024 -> 1024
BatchNormalization
HardTanh
FC_3
  1024 -> 10
BatchNormalization
```

<https://github.com/brycexu/BinarizedNeuralNetwork/tree/master/SecondTry>

The problems about the current BNN model ?

⇒ Accuracy Loss ! ⇒ Possible reasons ?

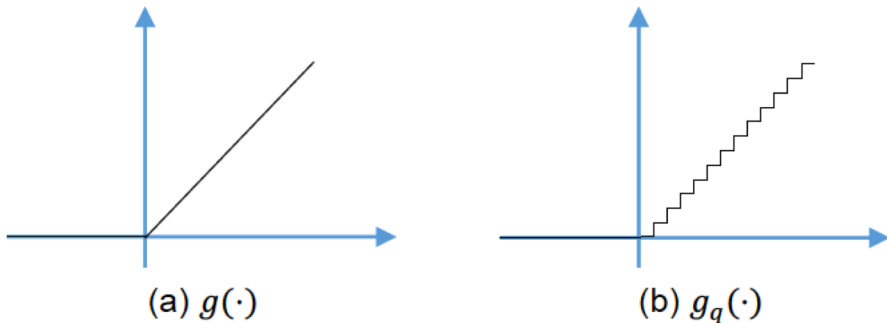
Problems 1: Robustness Issue

BNN always has larger output change which makes them more susceptible to input perturbation.

Problems 2: Stability Issue

BNN is hard to optimize due to problems such as gradient mismatch. This is because of the non-smoothness of the whole architecture.

Gradient mismatch:



The effective activation function in a fixed point network is a non-differentiable function in a discrete point network
That is why we cannot apply ReLU in BNN !

The potential ways to optimize BNN model ?

Robustness Issue

1. Adding more bits ?

-- Ternary model (-1,0,+1)

-- Quantization

Research shows that having more bits at activations improve model' robustness.

3. Adding more weights ?

-- WRPN

Stability Issue ?

1. Better activation function ?

2. Better back-propagation methods ?

2. Weakening learning rate ?

Research shows that higher learning rate can cause turbulence inside the model, so BNN needs finer tuning.

4. Modifying the architecture ?

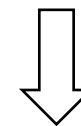
-- AdaBoost (BENN)

-- Recursively using binarization

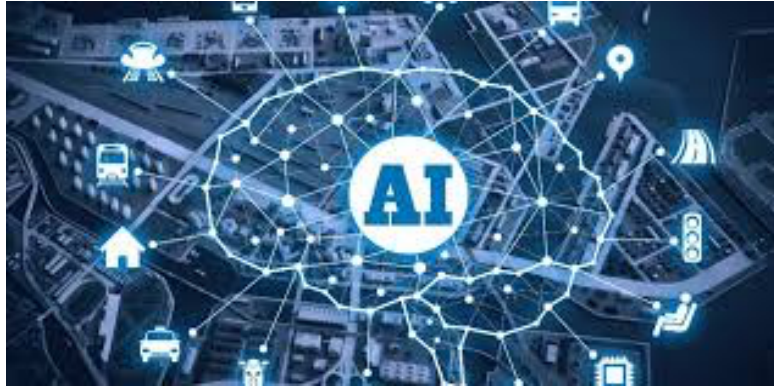
Table 4: Comparison with state-of-the-arts on ImageNet using AlexNet (W-weights, A-activation)

Method	W	A	Top-1
Full-Precision[34, 45]	32	32	56.6%
XNOR-Net [45]	1	1	44.2%
DoReFa-Net[58]	1	1	43.6%
BinaryConnect[10, 45]	1	32	35.4%
BNN[27, 45]	1	1	27.9%
BENN-Bag-5 (ours)	1	1	54.56%
BENN-Boost-5 (ours)	1	1	57.28%
BENN-Bag-8 (ours)	1	1	55.81%
BENN-Boost-8 (ours)	1	1	58.34%

More bits per network ?



More networks per bit ?



Thank you !

xxu373@wisc.edu